

pgstats in PostgreSQL 18

PostgreSQL Conference Japan 2025
November 21th, Tokyo
Michael Paquier

The lecturer

- French, based in Tokyo.
- PostgreSQL contributor since 2009
 - Committer
 - Hacker, Blogger, CF manager
- Website: <https://paquier.xyz>

Agenda

- Architecture
- Postgres 18
- Extensions

Architecture

Monitoring

- Aggregated view of cluster state
 - Objects: table, index, function
 - Processes: checkpointer, WAL receiver/sender, bgwriter
 - OS activity: WAL, I/O
 - Misc: Transaction, SLRU
- System views pg_stat_* (mostly, not pg_stat_activity!)
- <https://www.postgresql.org/docs/devel/monitoring-stats.html>

Up to Postgres 14

- Process: stats collector
- UDP for communication – pgstat_send()
and stats socket
- Stats on disk – Point of consistency
 - Written periodically, couple of times per second.
 - Loaded from files by each process
 - One global file, plus one file per database
- Performance bottleneck
 - Addition of new tenants
 - Multi-tenant, many files

Postgres 15 – Shared memory

- No stats collector.
- No periodic writes.
- Single on-disk file.
- Each stats kind is assigned an ID
- Commit 5891c7a8a.

Durability

- Single file
- Loaded by startup process
- Written/synced by last process at shutdown:
 - Checkpointer
 - Postmaster (single-user mode)
- Lost on crash
 - Autovacuum and table problem.
 - Relation stats in relfilenodes – recovery (Postgres 19)

Source code

- Code parts
 - pgstat.c
 - pgstat_shmem.c
 - pgstat_internal.h and pgstat.h
- One file for each stats kind:
 - pgstat_database.c
 - pgstat_archiver.c, etc.
- src/backend/utils/activity/

Fixed-sized stats

- Chunk of shared memory
- Memory allocated at startup
- LWLocks included in each stats structure
 - pgstat_internal.h
 - PgStatShared_*
- Counters cover a known size
 - WAL
 - I/O
 - SLRU

Variable-sized stats - Details

- Hash table in dynamic shared memory (dhash.c, DSA)
- Hash key:
 - Stats kind (4 bytes)
 - Database OID (4 bytes)
 - 8-byte ID (4-byte OID in v15~17)
- Objects: database, table, functions, DDL-related.
- Locking internal, partition of dhash.
- Maximum number not capped.

Variable-sized stats - Reporting

- Limit shared memory interaction.
- Pending stats
 - Get reference to entry in shmem, report pending activity.
 - Static to each process.
 - Saved in memory context created in Top MemoryContext.
- Flush by pgstat_report_stat()
 - Transaction commit.
 - Hardcoded max interval (60s).
 - Can be forced, waiting on locks.

Snapshots - stats_fetch_consistency

- Behavior of stats in a single transaction
 - “none”, each access fetches data from shmem.
 - “cache”, first entry access cached until end of transaction.
 - “snapshot”, caches all stats of database connected to.
- `pg_stat_clear_snapshot()` to clean up cache.

Postgres 18

Backend statistics

- Hash key based on proc number.
- Variable-sized.
- pgstat_backend.c
- For WAL and I/O, same fields as pg_stat_wal and pg_stat_io.
- Functions
 - pg_stat_get_backend_wal(pid), single row
 - pg_stat_get_backend_io(pid), multiple rows
 - pg_stat_reset_backend_stats()
 - Join with pg_stat_activity

pg_stat_io

- Multiple rows:
 - Type of operation (IOOp): write, read, syncs..
 - Context (IOContext): VACUUM, bulk-read, bulk-write, init (new in 18!)
 - Backend type: backend, checkpointer, autovacuum..
- Counters
 - In bytes in v18
 - BLCKSZ and number of operations up to v17.
- Data of pg_stat_wal moved to pg_stat_io

`pg_stat_database`

- Parallel worker activity
- Fields
 - `parallel_workers_to_launch`, number decided by planner.
 - `parallel_workers_launched`, number actually launched.
- Tuning of parallel workers

`pg_stat_statements` - Fields

- Like `pg_stat_database`, for each query.
 - Number of workers planned.
 - Number actually launched.
 - Available only in EXPLAIN and logs up to v17.
- `wal_buffers_full`
 - Count when WAL buffers become full
 - Useful for tuning of GUC `wal_buffers`

pg_stat_statements - Normalization

- IN/ANY clauses
 - SELECT * FROM tab WHERE a IN (\$1)
 - ORMs with large number of values
 - JDBC
- Relation name, not OID in query jumbling
 - Temp tables with same name: different query ID
 - Same table name + different schema: same query ID
 - Use alias in FROM clause for different query ID:
SELECT * FROM s1.tab AS s1tab;
SELECT * FROM s2.tab AS s2tab;

Extensions

Custom Cumulative Statistics

- Plugins and extensions can use pgstats!
- Similar to custom WAL RMGRs
 - Load with `shared_preload_libraries`
 - When removed from `shared_preload_libraries`, same as corrupted.
- Assigned fixed ID
- <https://www.postgresql.org/docs/develop/xfunc-c.html#XFUNC-ADDIN-CUSTOM-CUMULATIVE-STATISTICS>

Stats kind ID

- Whole range
 - PGSTAT_KIND_MIN 1
 - PGSTAT_KIND_MAX 32
- Invalid = 0
- Built-in: 1 to 12, up to 23 available
- Custom:
 - PGSTAT_KIND_CUSTOM_MIN = 24
 - PGSTAT_KIND_CUSTOM_MAX = PGSTAT_KIND_MAX
- pgstat_kind.h

Source code - Register

- PGSTAT_KIND_EXPERIMENTAL for development.
- ID overlap => stats file corruption.
- Reserve and ID:
<https://wiki.postgresql.org/wiki/CustomCumulativeStats>
- See pgstat_register_kind().

Callbacks - Basics

- PgStat_KindInfo in pgstat_internal.h
- Some boolean fields:
 - fixed_amount: variable or fixed size?
 - accessed_across_databases
 - write_to_file
- Some functions
 - flush_pending_cb => Variable-sized
 - flush_static_cb => Fixed-sized
 - init_backend_cb => Initialization action for backends

Callbacks – Names and Object ID mapping

- For on-disk storage
 - Define custom object IDs when reading from disk, based on names.
 - Translate object IDs to names, when writing.
- `to_serialized_name`
- `from_serialized_name`
- Example: replication slot stats

Source code - Templates

- Statistics for injection points
- Hash key:
 - Invalid database OID
 - Hash of injection point name.
- `src/test/modules/injection_points/`
 - `injection_stats.c` for variable-sized stats
 - `injection_stats_fixed.c` for fixed-sized stats
- GUC to control if stats are enabled.

`pg_stat_statements`?

- Query text file
 - Performance bottleneck
 - Move to DSA
- Deallocation and `pg_stat_statements.max`
 - Track number of entries in hash table of pgstats.
 - New callback for deallocation?
- Do NOT move in core. Move to custom pgstats. Try at least. Promise.

Thanks!
Questions?